

CPSC 261 – BASICS OF COMPUTER SYSTEMS

2017 Winter Term 2

Module 1: Memory models, C, pointers and assembly

1

LEARNING GOALS

- At the end of this module, you will be able to:
 - Explain the purpose of simple functions written in x86-64 assembly language.
 - Explain how a C program uses memory.
 - Write simple C functions that allocate and deallocate memory dynamically.
 - Recognize and fix common errors related to dynamic memory allocation.

MODULE SUMMARY

- **Module Summary**
 - The x86-64 assembly language
 - C and memory
 - Addresses and pointers
 - Dynamic memory allocation
 - Common programming mistakes with pointers
 - How a process' memory is organized
 - A look at the runtime stack
 - Implementing a dynamic memory allocator

AN ASIDE: LAB INFORMATION

- Linux is a free Unix-like operating system
 - Unix was invented at Bell Labs starting in 1969.
 - By 1973 it was in active use.
 - The original source code was encumbered with licensing concerns.
 - Linux is a re-implementation of the ideas.

AN ASIDE: IN-CLASS CODE

- All longer examples used in class will be available from a Stash repo.
- To create a local copy of the repo, use the command:
 - `git clone`
`https://stash.ugrad.cs.ubc.ca:8443/git/CS261_2017W2/examples.git`
- It will create a directory `examples` with the examples released so far (each in its own subdirectory).
- To update the repo when new examples are added to it, change to directory `examples`, and then use
 - `git pull`

THE X86-64 ASSEMBLY LANGUAGE

- We will use the x86-64
 - 16 general purpose registers (1 stack pointer, 1 frame pointer when needed).
 - instructions: mov, add, sub and lots more.
 - data sizes of 1, 2, 4, 8 (and occasionally 16) bytes.
 - parameters in registers unless there are too many (by convention).

THE X86-64 ASSEMBLY LANGUAGE

- How do you learn a new assembly language?
 - Read a short description (32 bits x86 only)
 - <http://www.cs.virginia.edu/~evans/cs216/guides/x86.html>
 - Read a longer description
 - <http://csapp.cs.cmu.edu/public/1e/public/docs/asm64-handout.pdf>
 - <https://software.intel.com/en-us/node/181178>
 - Or use a C compiler and read the assembly language code it produces.

X86-64 INSTRUCTIONS

- **Instruction types:**
 - Data movement (mov)
 - Arithmetic and logical (add, sub, mul, div, inc, and, or, not)
 - Comparisons (test, cmp)
 - Branches (jXX, jmp, loop)
 - Function call related (push, pop, call, ret)
 - Privileged (iret, hlt, cli, sti)
 - Will be covered in more detail later in the course

X86-64 INSTRUCTIONS (CONT.)

- The instruction's last letter indicates operands' size
 - b = 8 bits, l = 32 bits, q = 64 bits for integers
 - s = 32 bits, d = 64 bits, t = 128 bits for floating points
- The register's first letter indicates its size
 - ah, al: 8 bits
 - ax: 16 bits
 - eax: 32 bits
 - rax: 64 bits
- Not all sizes are available for all registers.

X86-64 ADDRESSING MODES

- Addressing modes:
 - register direct: `%rax`
 - indirect: `offset(%rax)`
 - indexed: `offset(%rax, %rbx)`
 - indexed scaled: `offset(%rax, %rbx, 4)`
 - absolute: `address`
 - immediate: `$value`

ASSEMBLY CONVENTIONS

- There are conventions:
 - What registers are used for what?
 - How are values passed to functions?
 - How are values returned from functions?
 - What if a value is too large to fit in a register?
- Conventions may vary by:
 - CPU architecture
 - Operating system
 - Compiler

X86-64 CALLING CONVENTIONS

- Stack pointer: `%rsp`
- Arguments passed in registers, in this order: `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8`, `%r9`
 - Additional arguments passed in stack
- Returning value passed in `%rax`
 - Use `%rax` and `%rdx` if returning value is 128 bits
 - Use memory for larger values
- Caller-save registers (caller's responsibility to save before calling other functions): `%rax`, `%r10`, `%r11` + arguments
- Callee-save registers (function must restore value before returning if changed): `%rbx`, `%rbp`, `%r12`, `%r13`, `%r14`
 - In x86-64: `%r15`*

SOME EXAMPLES

- Let us look at examples (in the examples repo):
 - odd.c
 - sum.c

REVIEW: ADDRESSES AND POINTERS

- Consider the following (incomplete) piece of code:

```
int *a, b;  
int **c;
```

```
...
```

```
d = &c;  
x = a + 7;  
y = c[3];
```

- Describe the type of value stored in each variable without using the word “pointer”.

REVIEW: ADDRESSES AND POINTERS

- What does the following C function print?

```
void do_something() {  
    char mtl[] = "      ";  
    char *qc, **cdn = &qc;  
  
    mtl[0] = 'Y';  
    mtl[4] = '\\0';  
    qc = mtl + 2;  
    *qc = 'A';  
    qc[-1] = 'P';  
    *cdn = qc + 1;  
    **cdn = 'Z';  
    *(qc - 1) = 'E';  
    printf("%s\\n", mtl);  
}
```

REVIEW: ADDRESSES AND POINTERS

- Rewrite the following piece of code using arrays to make it more readable.

```
int confusing(long *p, int n)
{
    long *q = p + n;
    while (n > 0 && *p++ == *--q)
    {
        n -= 2;
    }
    return n <= 0;
}
```


LOCAL VARIABLES AND ARGUMENTS

```
void b (int a0, int a1) {  
    int l0 = 0;  
    int l1 = 1;  
}
```

- Can `l0`, `l1`, `a0`, `a1` be allocated statically?
 - In other words, can their static address be determined at compilation time?
 - What about recursion?

LOCAL VARIABLES AND RECURSION

- How many different versions of `n` are there?

```
void a (int n) {  
    if (n == 0) return;  
    a(n - 1);  
    printf("%d\n", n);  
}
```

- What if there is no apparent recursion?

```
void b (int n) {  
    int l = n * n;  
    c(l - 1);  
}
```

- What if `c()` calls `b()`?

LIFE OF A LOCAL VARIABLE / ARGUMENT

- **Scope**
 - Local variables are only accessible within declaring procedure
 - Each execution has its own private copy
- **Lifetime**
 - Allocated when procedure starts
 - “Freed” when procedure returns (in most languages)

PROCEDURE ACTIVATION

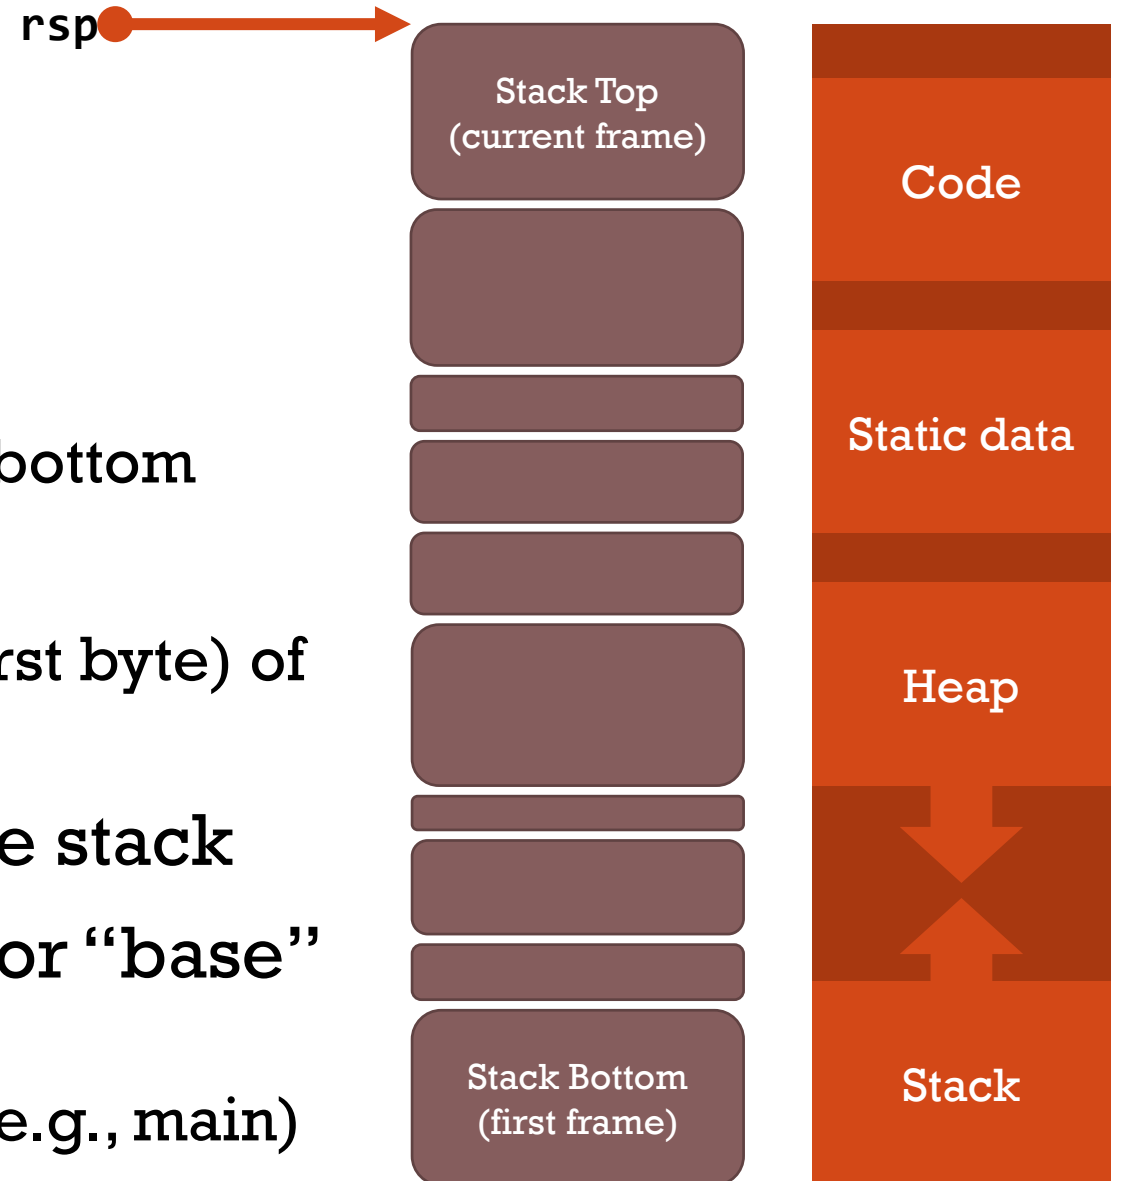
- **Activation: execution of a procedure**
 - Starts with procedure is called, ands when it returns
 - There can be many activations of same procedure alive at once
- **Activation Frame**
 - memory that stores activation's state
 - Includes local variables and arguments
- **Should we allocate activation frames from the heap?**
 - Call `malloc()` to create frame on procedure call?
 - Call `free()` on procedure return?

HEAP VS ALLOCATION FRAMES

- Order of frame allocation and deallocation is special
 - freed in reverse order of allocation
- Simple allocation for frames:
 - Reserve big chunk of memory for all frames
 - Initial address known
 - Simple, cheap allocation: add or subtract from a pointer
- Questions
 - What data structure is this like?
 - What restriction do we place on lifetime of local variables?

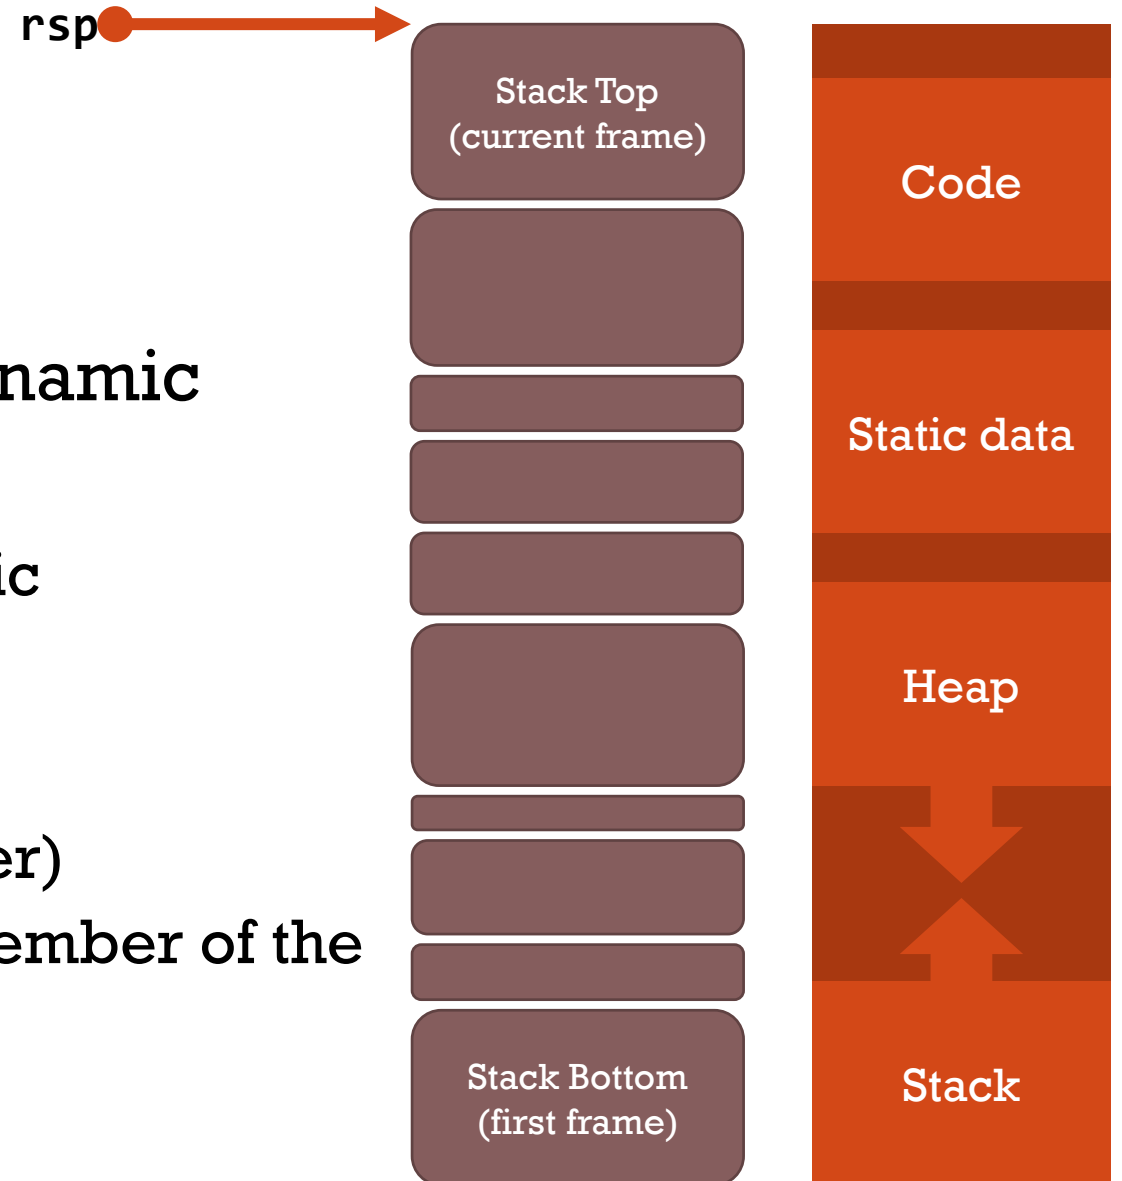
RUNTIME STACK

- Stack of activation frames
 - Stored in memory, grows up from bottom
- Stack pointer
 - Stores base address (address of first byte) of current frame
- Current frame is the “top” of the stack
- First activation is the “bottom” or “base” of the stack
 - Local variables of initial function (e.g., main)



VARIABLE ADDRESSES

- Value of the stack pointer is dynamic
- Local variables and arguments
 - Size of each frame is (usually) static
 - Offset from stack pointer is static
- Each frame is like a struct
 - Top of frame is in `rsp` (stack pointer)
 - Each variable in procedure is a member of the struct



WHAT IS STORED IN THE ACTIVATION FRAME?

- Local variables
- Arguments
 - Some architectures use registers for arguments
- Return address
- Other saved registers
 - Called function may change register values
 - Values that must be kept after call are saved

saved registers...
local variables...
return address
arguments...

saved registers...

SOME IMPLICATIONS

- What is the value of `l` in `foo` when it is active?

```
void goo() { int x = 3; }           goo();  
void foo() { int l; }              foo();
```

- What is wrong with this?

```
int *foo() {  
    int l;  
    return &l;  
}
```

PROCESS MEMORY ORGANIZATION

- When a program is executed:
 - Space is allocated for the shared libraries it needs
 - Instructions and initialized data are loaded in memory
 - Space is reserved for uninitialized data
- The stack and the heap are set up
 - The stack is managed by the compiler's code
 - The heap is managed by the user's program
- On a Linux system, we can look at `/proc/pid/maps` to see how memory is used for process *pid*.

PROCESS MEMORY EXAMPLE

00400000-00401000	r-xp	00000000	08:08	5244791	/home/patrice/fib
00600000-00601000	r--p	00000000	08:08	5244791	/home/patrice/fib
00601000-00602000	rw-p	00001000	08:08	5244791	/home/patrice/fib
7f86a3122000-7f86a32dd000	r-xp	00000000	08:02	1308509	/lib/x86-64-linux-gnu/libc-2.19.so
7f86a32dd000-7f86a34dd000	---p	001bb000	08:02	1308509	/lib/x86-64-linux-gnu/libc-2.19.so
7f86a34dd000-7f86a34e1000	r--p	001bb000	08:02	1308509	/lib/x86-64-linux-gnu/libc-2.19.so
7f86a34e1000-7f86a34e3000	rw-p	001bf000	08:02	1308509	/lib/x86-64-linux-gnu/libc-2.19.so
7f86a34e3000-7f86a34e8000	rw-p	00000000	00:00	0	
7f86a34e8000-7f86a35ed000	r-xp	00000000	08:02	1308538	/lib/x86-64-linux-gnu/libm-2.19.so
7f86a35ed000-7f86a37ec000	---p	00105000	08:02	1308538	/lib/x86-64-linux-gnu/libm-2.19.so
7f86a37ec000-7f86a37ed000	r--p	00104000	08:02	1308538	/lib/x86-64-linux-gnu/libm-2.19.so
7f86a37ed000-7f86a37ee000	rw-p	00105000	08:02	1308538	/lib/x86-64-linux-gnu/libm-2.19.so
7f86a37ee000-7f86a3811000	r-xp	00000000	08:02	1308526	/lib/x86-64-linux-gnu/ld-2.19.so
7f86a39e5000-7f86a39e8000	rw-p	00000000	00:00	0	
7f86a3a0e000-7f86a3a10000	rw-p	00000000	00:00	0	
7f86a3a10000-7f86a3a11000	r--p	00022000	08:02	1308526	/lib/x86-64-linux-gnu/ld-2.19.so
7f86a3a11000-7f86a3a12000	rw-p	00023000	08:02	1308526	/lib/x86-64-linux-gnu/ld-2.19.so
7f86a3a12000-7f86a3a13000	rw-p	00000000	00:00	0	
7ffffc88a3000-7ffffc88c4000	rw-p	00000000	00:00	0	[stack]
7ffffc88cf000-7ffffc88d1000	r-xp	00000000	00:00	0	[vdso]
ffffffffffff600000-ffffffffffff601000	r-xp	00000000	00:00	0	[vsyscall]

DYNAMIC MEMORY ALLOCATION

- Allocating memory
 - We use:
 - `void *malloc(size_t size);`
 - `void *calloc(size_t count, size_t size);`
 - Typical usage:
 - Given
 - `type *ptr;`
 - To allocate memory for one object:
 - `ptr = malloc(sizeof(type));`
 - To allocate memory for an array of objects:
 - `ptr = calloc(count, sizeof(type));`

DYNAMIC MEMORY ALLOCATION

- Deallocating memory:
 - We use:
 - `void free(void *ptr);`
 - Example:
 - `free(ptr_to_type);`
 - This works for a single object, as well as for arrays

DYNAMIC MEMORY ALLOCATION

- Example: allocate and return a string with the English alphabet.

```
char *create_alphabet_string()  
{
```

```
}
```

DYNAMIC MEMORY ALLOCATION

```
char *create_alphabet_string() {  
    char *my_string = malloc(27 * sizeof(char));  
    int i;  
    for (i = 0; i < 26; i++) {  
        my_string[i] = 'a' + i;  
    }  
    my_string[26] = '\0';  
    return my_string;  
}
```

DYNAMIC ALLOCATION: USAGE

- Example: array that grows as data grows
 - When adding an element, if array is full, allocate bigger space and copy old data to new space
- Example: linked list
 - Each element is a struct that contains a pointer to next element
 - Lab assignment #2

DANGLING POINTERS

- Dangling pointer: pointer that does not point to actual usable memory
- Examples:
 - Uninitialized pointer
 - Multiple pointers to same location, one is freed and the other is still in use
 - Calling free on the same memory twice
 - Function returns pointer to local variable
- Good practices to avoid dangling pointers:
 - Initialize all pointers to valid data (e.g., NULL)
 - If needed, implement reference counting

MEMORY LEAKS

- Memory leak: memory has been allocated but not deallocated
 - Usually caused by lost reference to dynamically allocated data
- Examples:
 - Function allocates memory, then returns without saving or returning memory
 - Function returns dynamically allocated memory, but return value is ignored
 - Last pointer to allocated space is changed to different value

AVOIDING MEMORY LEAKS

- Good programming practices to avoid leaks:
 - When possible, allocate and free data in same function
 - Check old value of a pointer before changing it
 - Point it to NULL if not assigned to anything
 - If needed, implement reference counting

OTHER COMMON MISTAKES WITH POINTERS

- Buffer overflow: using more data than allocated
 - Can be a problem with global and local arrays as well
 - Can be caused by off-by-one errors (e.g., not counting the string termination byte)

IMPLEMENTING A DYNAMIC MEMORY ALLOCATOR

- So far we have been using malloc/calloc and free
- But how are they implemented?
 - How is available memory maintained?
 - How to keep track of freed blocks?
 - When is it ok to reuse a block?

MEMORY ALLOCATOR REQUIREMENTS

- Handling arbitrary sequences of requests
 - The allocator can not control the requests for malloc/free
- Making immediate responses to requests
 - The allocator can not wait to process several requests at once, even if it would be more efficient
- Using only the heap for its data structures
 - We can use a constant amount of additional space only
- Not modifying allocated blocks
 - The user program assumes their contents won't change
 - It also assumes its location won't change

MEMORY ALLOCATOR IMPLEMENTATION

- Implementation issues:
 - *Placement*: when malloc() is called, how to we find a free block that will be used to satisfy the request?
 - *Splitting*: if we only need part of a free block to satisfy a request, what do we do with the rest?
 - *Coalescing*: do we merge a newly free()'d block with adjacent free blocks?
- These issues arise in all implementations

MEMORY ALLOCATOR IMPLEMENTATION

- **Goals:**
 - Maximizing throughput
 - We want to respond to requests quickly
 - So we need to use simple data structures
 - Maximizing memory utilization
 - We want to avoid internal and external fragmentation

FRAGMENTATION


- Internal fragmentation
 - Most allocators impose a minimum size on the blocks they return to a process
 - Because of alignment requirements (pointer addresses must be a multiple of 4, or 8)
 - Because the allocation needs to store information inside the blocks once they are freed; so the blocks must be large enough
 - Hence a request for a very small amounts of memory returns a larger block than that requested
 - Some space inside the block is wasted

FRAGMENTATION


- External fragmentation
 - Unlike files on disk, the block returned by `malloc()` must consist of consecutive memory locations
 - Sometimes there may be enough free space in total, but no free block is large enough to satisfy the request
 - The more calls to `malloc()` and `free()` have been made with requests for different size blocks, the more external fragmentation is a problem

IMPLICIT FREE LIST

- A simple implementation: the implicit free list.



- We have a linked list of blocks.
 - The list contains both occupied and free blocks.
 - Each block contains its size.
 - Each block knows if it's free or occupied:




IMPLICIT FREE LIST: SPLITTING

- If the block used is larger than the requested size, we can
 - use the whole block (increases internal fragmentation)
 - divide the block in two (may end up with many small blocks)


IMPLICIT FREE LIST: COALESCING

- When freeing a block, merge adjacent free blocks
- Avoids ending up with lots of small free blocks all adjacent
- We can do this on every free
 - Advantage: simpler
 - Disadvantage: free operation becomes slower
 - Alternative: wait until an allocation request fails
- Problem: how to find adjacent blocks



IMPLICIT FREE LIST: COALESCING

- We store the block size at the end of the block also:



IMPLICIT FREE LIST: PLACEMENT

- Placement:

- *First-fit*: return the first free block that is large enough
 - retains large free block near end of the list
 - Disadvantage: search time if too many small blocks
- *Next-fit*: similar but start searching from the last allocated block
 - Advantage: faster search time
 - Disadvantage: worse memory utilization than first-fit
- *Best-fit*: find the free block whose size is closest to the requested size
 - Advantage: optimal use of memory
 - Disadvantage: slower

EXPLICIT FREE LIST

- *Explicit free list*: uses the payload in free blocks to point to other free blocks
 - Block search is faster (don't need to check blocks in use)
 - Doubly-linked list
 - Disadvantage: minimum payload size must be enough for two pointers
- Linked-list order (where to pointers point to)
 - Last-in First-out: free blocks go at the beginning of the list
 - `free()` takes constant time
 - In address order: pointers list blocks in address order
 - `free()` requires linear time (must search previous/next free blocks)
 - We get slightly better memory utilization (search is in memory order)

SEGREGATED FREE LIST

- *Segregated free list*: one linked list per block size
 - Blocks point to other blocks of similar size
- One approach (segregated fits):
 - malloc():
 - find a block large enough,
 - split it if desired, and insert the other piece in the appropriate free list.
 - free():
 - coalesce with adjacent free blocks if possible,
 - store the new free block in the appropriate free list.

SEGREGATED FREE LIST (CONT.)

- Advantages of segregated free lists:
 - Searching for free blocks is more efficient
 - We are only searching part of the heap
 - Memory utilization improves
 - First-fit search with a segregated list approximates a best-fit search of the entire heap
- The GNU malloc package, part of the standard C library on all Linux systems, uses segregated free lists